

CREDIT CARD FRAUD DETECTION

- **Credit card fraud detection is a process and set of**
- **techniques used by financial institutions, merchants, and**
- **credit card companies to identify and prevent**
- **unauthorized or fraudulent use of credit cards. It involves**
- **the use of various methods to analyze and monitor**
- **credit card transactions in real-time to detect any**
- **unusual or suspicious activity. Some common:**



PROBLEM STATEMENT

- **DEVELOP AN EFFECTIVE AND SCALABLE CREDIT CARD FRAUD DETECTION SYSTEM USING MACHINE LEARNING TECHNIQUES TO IDENTIFY AND PREVENT FRAUDULENT TRANSACTIONS IN REAL-TIME, MINIMIZING FINANCIAL LOSSES AND ENSURING A SEAMLESS USER EXPERIENCE FOR LEGITIMATE CUSTOMERS.**

DESIGN THINKING

- DESIGN THINKING IS A USER-CENTERED APPROACH TO PROBLEM-SOLVING THAT INVOLVES EMPATHIZING WITH USERS, DEFINING THE PROBLEM, IDEATING SOLUTIONS, PROTOTYPING, AND TESTING.
- PHASES OF DESIGN THINKING IN CREDIT CARD FRAUD DETECTION:
 - 1. EMPATHIZE
 - 2. DEFINE
 - 3. IDEATE
 - 4. PROTOTYPE
 - 5. TEST.

EMPATHIZING WITH USER

- UNDERSTAND THE NEEDS AND PAIN POINTS OF BOTH CUSTOMERS AND FINANCIAL INSTITUTIONS.
- COLLECT AND ANALYZE DATA ON PAST FRAUDULENT TRANSACTIONS AND THEIR IMPACT ON CUSTOMERS AND BUSINESSES.
- INTERVIEW FRAUD ANALYSTS, SECURITY EXPERTS, AND CUSTOMERS TO GAIN INSIGHTS INTO THEIR PERSPECTIVES AND CHALLENGES.

DEFINING THE PROBLEM:

- **Clearly define the problem by identifying the main issues in credit card fraud detection.**
- **Create a problem statement that addresses the key concerns, such as reducing false positives, improving detection accuracy, and enhancing customer experience.**


IDEATING

SOLUTIONS:




- Developing advanced anomaly detection algorithms.
 - Implementing real-time transaction monitoring.
 - Enhancing user authentication and verification processes.
 - Exploring behavior-based modeling for fraud detection.
 - Using deep learning models for pattern recognition.
-
- Encourage cross-functional collaboration among data scientists, engineers, domain experts, and UX designers to brainstorm innovative solutions.

PROTOTYPE



- **Create prototypes or mock-ups of potential solutions. For instance:**
- **Develop a user interface for fraud analysts to investigate suspicious transactions efficiently.**
 - ▶ **Create a machine learning model prototype for fraud detection and prevention.**
- **These prototypes should be low-cost and**

TESTING



- Gather feedback from stakeholders, including fraud analysts, customers, and technical experts, on the prototypes.
- Iterate on the prototypes based on the feedback received.
- Conduct simulations or pilot tests to evaluate the effectiveness of the proposed solutions.

DATA SET COLLECTION

The first step is to collect reliable data so that your machine learning

model can find the correct patterns. The quality of the data that you feed to

the machine will determine how accurate your model

DATASET LINK:

<https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>

3.16E+09	500	26000 N	0 Y	Y	800	677.2	6 Y
3.16E+09	500	27000 N	0 Y	Y	800	677.2	6 Y
3.16E+09	262.5	11287.5 N	0 N	N	900	345.5	7 Y
3.162E+09	185.5	11130 Y	20 N	N	0	0	0 Y
3.162E+09	185.5	6121.5 Y	20 N	N	0	0	0 Y
3.162E+09	185.5	7049 Y	20 N	N	0	0	0 Y
3.356E+09	166.78847	4836.8657 N	0 N	N	721	229	9 Y
3.359E+09	444.99701	21804.854 N	0 Y	Y	0	0	0 Y
3.36E+09	152.45157	4116.1923 N	0 Y	Y	865	375	8 Y
3.365E+09	36.919488	2141.3303 N	5 Y	Y	0	0	0 Y
3.365E+09	806.17954	23379.207 N	0 N	N	816	811	5 Y
3.37E+09	257.09117	10283.647 N	4 Y	N	0	0	0 Y
3.376E+09	601.45297	24659.572 N	7 N	N	0	0	0 Y
3.387E+09	222.52982	12461.67 N	9 Y	N	0	0	0 Y
3.388E+09	231.06732	12708.702 N	0 N	N	986	650	8 Y
3.395E+09	675.56948	39858.599 N	9 Y	N	0	0	0 Y
3.402E+09	242.15174	10654.677 N	0 Y	Y	0	0	0 Y
3.406E+09	804.76177	42652.374 N	0 N	N	953	950	8 Y
3.408E+09	432.03025	22033.543 N	8 Y	Y	0	0	0 Y
3.413E+09	356.81226	16056.552 N	9 Y	Y	0	0	0 Y
3.418E+09	456.28312	15057.343 N	0 Y	Y	0	0	0 Y
3.45E+09	780.6885	36692.359 N	0 Y	N	896	839	6 Y
3.462E+09	947.48903	44531.984 N	7 Y	N	0	0	0 Y
3.466E+09	172.51577	5002.9572 N	9 Y	Y	0	0	0 Y
3.484E+09	111.37507	5680.1286 N	2 Y	N	0	0	0 Y
3.48E+09	670.04033	38076.550 N	0 Y	Y	0	0	0 Y






DATA PREPROCESSING

clean and preprocess the
data to remove noise
and inconsistencies
handle missing values
and outliers






FEATURE ENGINEERING

- Create a relevant features from the data such as transaction, frequency, Location based features and user behavior patterns

MODEL SELECTION

- choose appropriate machine learning or deep learning models for fraud detection, such as logistic regression, decision trees, random forests, or neural networks. Consider ensemble methods for improved performance



TRAIN AND TEST

- To train and test a model for credit card fraud detection, you should follow these steps, which include splitting your dataset into training and testing sets:
- Data Preparation: Gather and preprocess your credit card transaction dataset. This dataset should include labeled data, where each transaction is classified as either legitimate or fraudulent.
- Data Splitting: Divide your dataset into three subsets: training, validation, and testing. A common split might be 60% for training, 20% for validation, and 20% for testing. The training set is used to train the model, the validation set helps tune hyperparameters, and the testing set is used for the final evaluation





- **Training Set:** The training set is used to train your machine learning or deep learning model. It should contain a diverse representation of both legitimate and fraudulent transactions to ensure the model learns to distinguish between the two classes.
- **Validation Set:** The validation set is used during the training process to fine-tune model hyperparameters and assess the model's performance. You can use it to prevent overfitting.

-
- **Testing Set:** The testing set is kept separate and used only after model training is complete. It is used to evaluate the model's performance on unseen data, simulating real-world scenarios.
 - **Model Training:** Train your chosen model using the training data, adjusting hyperparameters and architecture as needed. Common models for fraud detection include logistic regression, decision trees, random forests, and deep neural networks.



1. Dataset Information

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount, this feature can be used for example- dependant cost-sensitive learning. Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

Given the class imbalance ratio, we recommend measuring the accuracy using the Area Under the Precision-Recall Curve (AUPRC). Confusion matrix accuracy is not meaningful for unbalanced classification.

2. Import modules

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

0.3 Loading the dataset

```
[ ]: df = pd.read_csv("C:\Users\sdeva\Downloads\archive(1)\creditcard.csv")
df.head()
```

```
[ ]: 
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	


```
3  1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4  2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941
```

```
      V8      V9  ...      V21      V22      V23      V24      V25  \
0  0.098698  0.363787  ...-0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102 -0.255425  ...-0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ... 0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ...-0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739  ...-0.009431  0.798278 -0.137458  0.141267 -0.206010
```

```
      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053  149.62      0
1  0.125895 -0.008983  0.014724   2.69      0
2 -0.139097 -0.055353 -0.059752  378.66      0
3 -0.221929  0.062723  0.061458  123.50      0
4  0.502292  0.219422  0.215153   69.99      0
```

[5 rows x 31 columns]

```
[3]: # statistical info
df.describe()
```

```
[3]:
count      284807.000000  2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05  \
mean       94813.859575  3.918649e-15  5.682686e-16 -8.761736e-15  2.811118e-15
std        47488.145955  1.958696e+00  1.651309e+00  1.516255e+00  1.415869e+00
min          0.000000 -5.640751e+01 -7.271573e+01 -4.832559e+01 -5.683171e+00
25%        54201.500000 -9.203734e-01 -5.985499e-01 -8.903648e-01 -8.486401e-01
50%        84692.000000  1.810880e-02  6.548556e-02  1.798463e-01 -1.984653e-02
75%       139320.500000  1.315642e+00  8.037239e-01  1.027196e+00  7.433413e-01
max       172792.000000  2.454930e+00  2.205773e+01  9.382558e+00  1.687534e+01

count      2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05  \
mean     -1.552103e-15  2.040130e-15 -1.698953e-15 -1.893285e-16 -3.147640e-15
std        1.380247e+00  1.332271e+00  1.237094e+00  1.194353e+00  1.098632e+00
min     -1.137433e+02 -2.616051e+01 -4.355724e+01 -7.321672e+01 -1.343407e+01
25%     -6.915971e-01 -7.682956e-01 -5.540759e-01 -2.086297e-01 -6.430976e-01
50%     -5.433583e-02 -2.741871e-01  4.010308e-02  2.235804e-02 -5.142873e-02
75%        6.119264e-01  3.985649e-01  5.704361e-01  3.273459e-01  5.971390e-01
max        3.480167e+01  7.330163e+01  1.205895e+02  2.000721e+01  1.559499e+01

count      ...      V21      V22      V23      V24  \
mean      ...  1.473120e-16  8.042109e-16  5.282512e-16  4.456271e-15
std      ...  7.345240e-01  7.257016e-01  6.244603e-01  6.056471e-01
min      ... -3.483038e+01 -1.093314e+01 -4.480774e+01 -2.836627e+00
```

```
25%    ...-2.283949e-01 -5.423504e-01 -1.618463e-01 -3.545861e-01
50%    ...-2.945017e-02  6.781943e-03 -1.119293e-02  4.097606e-02
75%    ... 1.863772e-01  5.285536e-01  1.476421e-01  4.395266e-01
max     ... 2.720284e+01  1.050309e+01  2.252841e+01  4.584549e+00
```

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	1.426896e-15	1.701640e-15	-3.662252e-16	-1.217809e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 31 columns]

```
[4]: # datatype info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Time    284807 non-null  float64
1     V1      284807 non-null  float64
2     V2      284807 non-null  float64
3     V3      284807 non-null  float64
4     V4      284807 non-null  float64
5     V5      284807 non-null  float64
6     V6      284807 non-null  float64
7     V7      284807 non-null  float64
8     V8      284807 non-null  float64
9     V9      284807 non-null  float64
10    V10     284807 non-null  float64
11    V11     284807 non-null  float64
```

```
12 V12      284807 non-null float64
13 V13      284807 non-null float64
14 V14      284807 non-null float64
15 V15      284807 non-null float64
16 V16      284807 non-null float64
17 V17      284807 non-null float64
18 V18      284807 non-null float64
19 V19      284807 non-null float64
20 V20      284807 non-null float64
21 V21      284807 non-null float64
22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class    284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

0.4 Preprocessing the dataset

```
[5]: # check for null values
df.isnull().sum()
```


```
[5]: Time      0
V1          0
V2          0
V3          0
V4          0
V5          0
V6          0
V7          0
V8          0
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
```

```
V20      0
V21      0
V22      0
V23      0
V24      0
V25      0
V26      0
V27      0
V28      0
Amount    0
Class     0
dtype: int64
```

0.5 Exploratory Data Analysis

```
[6]: sns.countplot(df['Class'])
```

```
[6]: <AxesSubplot:xlabel='Class', ylabel='count'>
```

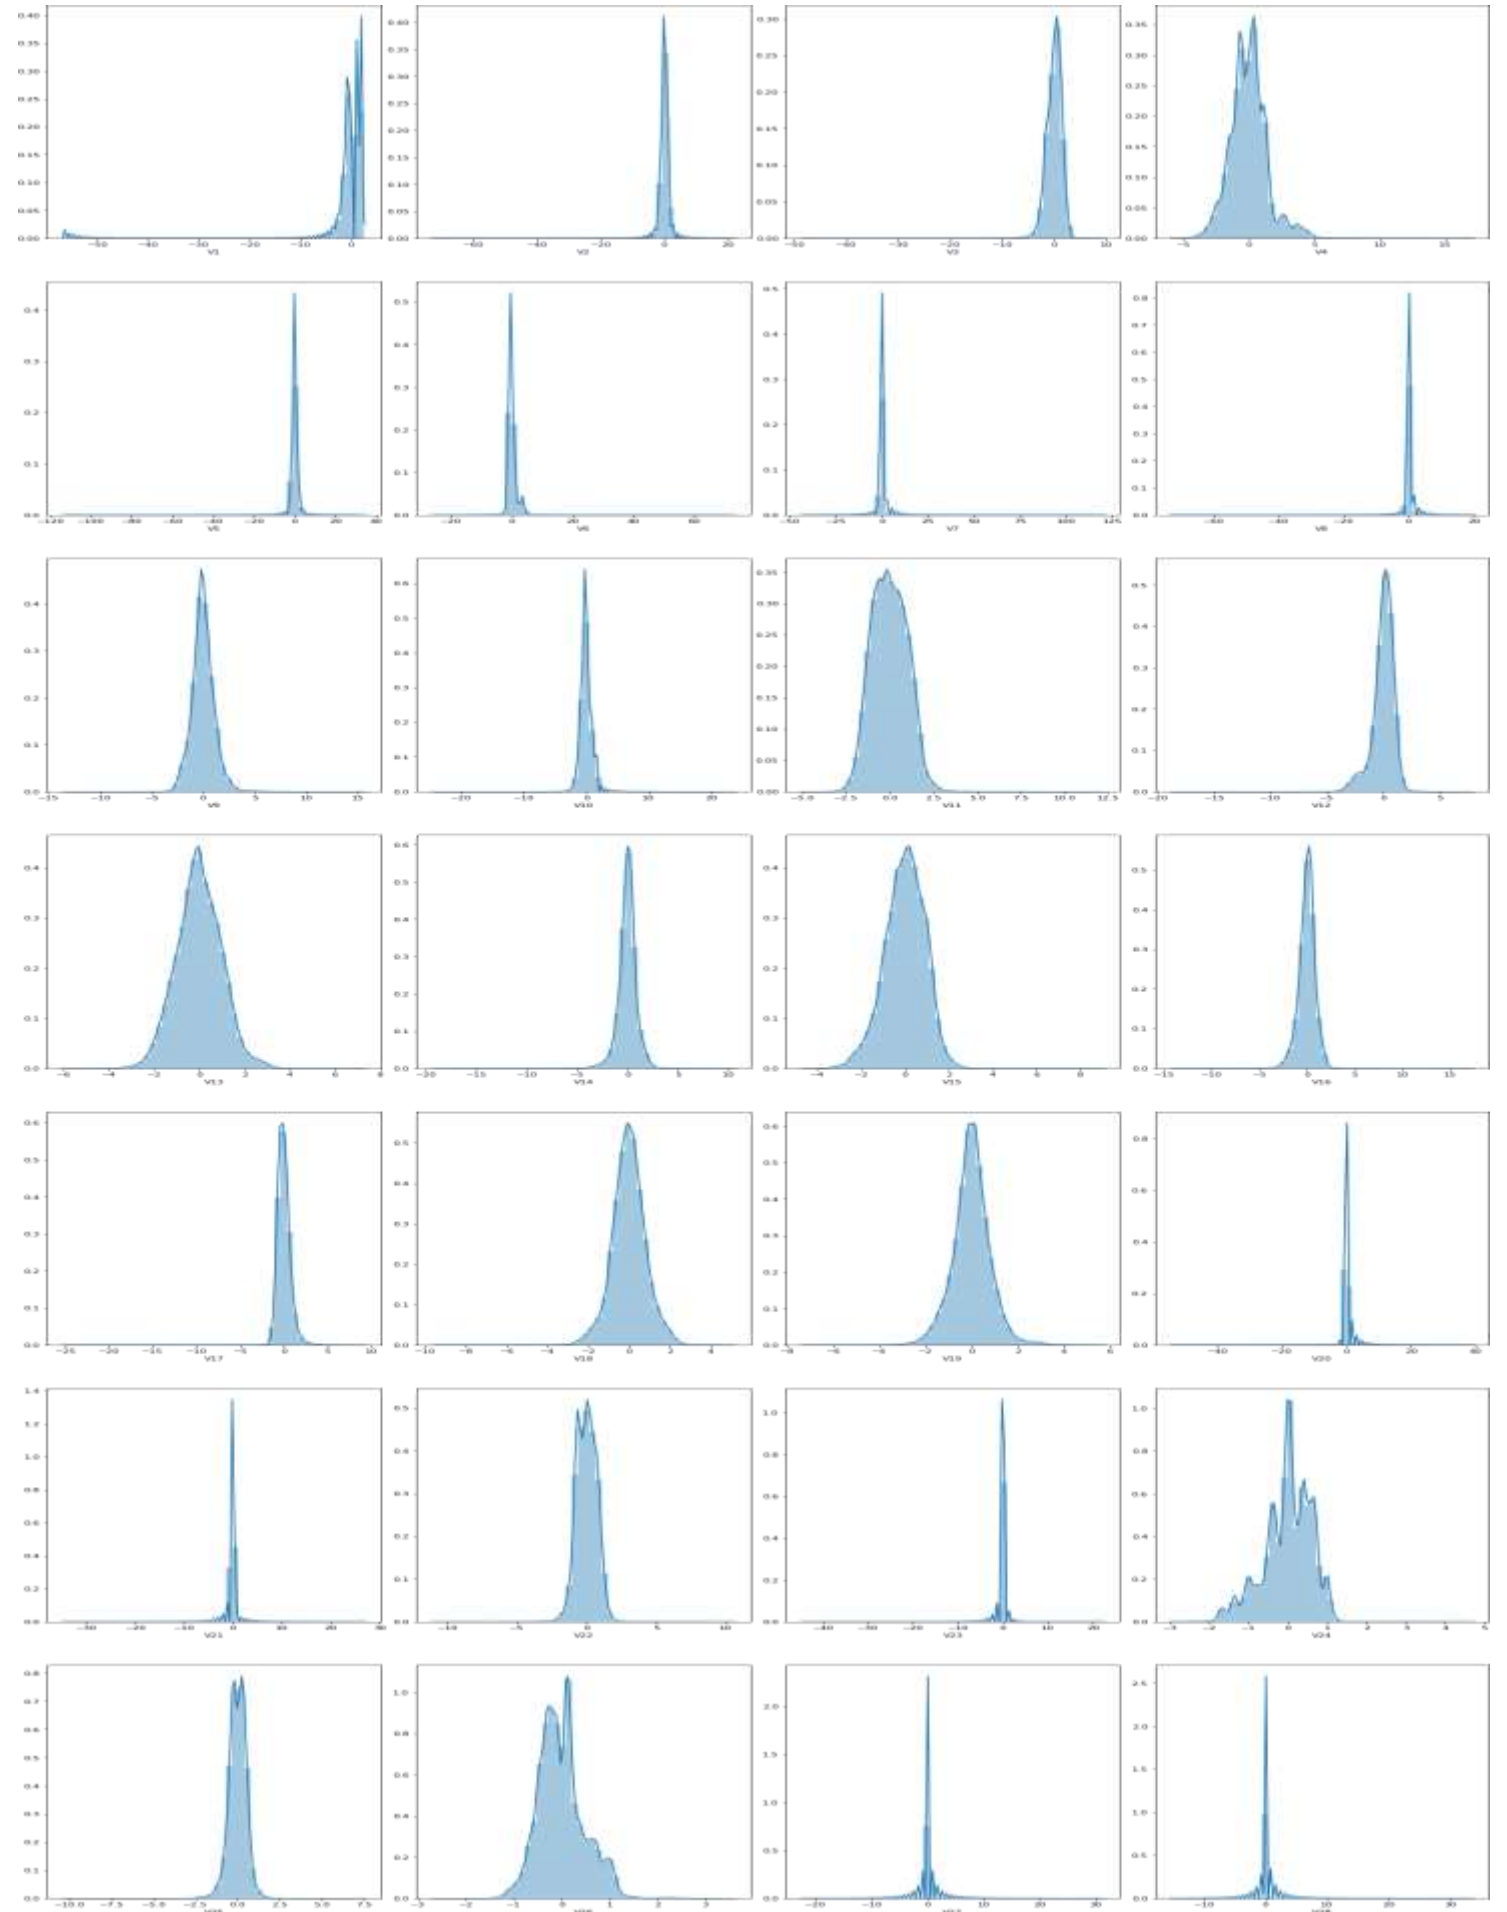


```
[10]: df_temp = df.drop(columns=['Time', 'Amount', 'Class'], axis=1)
```

```
# create dist plots
fig, ax = plt.subplots(ncols=4, nrows=7, figsize=(20, 50))
index = 0
```

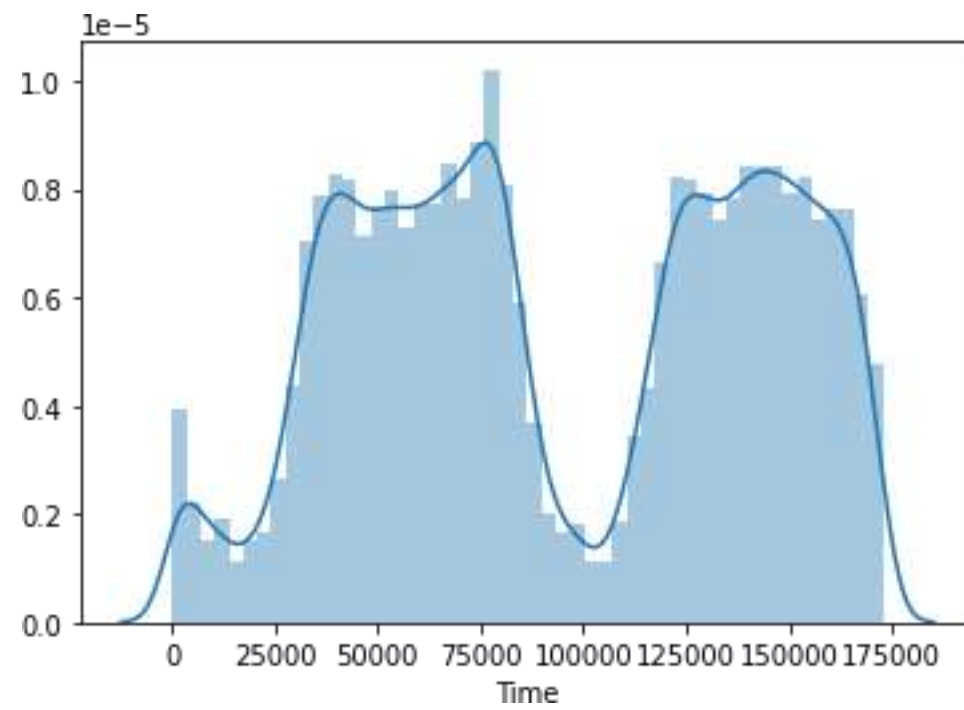
```
ax = ax.flatten()
```

```
for col in df_temp.columns:
    sns.distplot(df_temp[col], ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=5)
```




```
[11]: sns.distplot(df['Time'])
```

```
[11]: <AxesSubplot:xlabel='Time'>
```



```
[12]: sns.distplot(df['Amount'])
```


```
[12]: <AxesSubplot:xlabel='Amount'>
```



0.6 Coorelation Matrix

```
[14]: corr = df.corr()
plt.figure(figsize=(30,40))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

[14]: <AxesSubplot:>



0.7 Input Split

```
[15]: X = df.drop(columns=['Class'], axis=1)
      y = df['Class']
```

0.8 Standard Scaling

```
[16]: from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      x_scaler = sc.fit_transform(X)
```

```
[18]: x_scaler[-1]
```

```
[18]: array([ 1.64205773, -0.27233093, -0.11489898,  0.46386564, -0.35757    ,
            -0.00908946, -0.48760183,  1.27476937, -0.3471764 ,  0.44253246,
            -0.84072963, -1.01934641, -0.0315383 , -0.18898634, -0.08795849,
             0.04515766, -0.34535763, -0.77752147,  0.1997554 , -0.31462479,
             0.49673933,  0.35541083,  0.8861488 ,  0.6033653 ,  0.01452561,
            -0.90863123, -1.69685342, -0.00598394,  0.04134999,  0.51435531])
```

0.9 Model Training

```
[23]: # train test split
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report, f1_score
      x_train, x_test, y_train, y_test = train_test_split(x_scaler, y, test_size=0.
      ↪25, random_state=42, stratify=y)
```

```
[25]: from sklearn.linear_model import LogisticRegression
      model = LogisticRegression()
      # training
      model.fit(x_train, y_train)
      # testing
      y_pred = model.predict(x_test)
      print(classification_report(y_test, y_pred))
      print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.85	0.63	0.72	123
accuracy			1.00	71202
macro avg	0.92	0.81	0.86	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.719626168224299


```
[26]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
#training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", fl_score(y_test, y_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.95	0.76	0.85	123
accuracy			1.00	71202
macro avg	0.97	0.88	0.92	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.846846846846847

```
[37]: from xgboost import XGBClassifier
model = XGBClassifier(n_jobs=-1)
#training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", fl_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.94	0.80	0.86	123
accuracy			1.00	71202
macro avg	0.97	0.90	0.93	71202
weighted avg	1.00	1.00	1.00	71202


F1 Score: 0.8634361233480178



```
[29]: # hint - use combination of over sampling and under sampling  
# balance the class with equal distribution  
from imblearn.over_sampling import SMOTE  
over_sample = SMOTE()  
x_smote, y_smote = over_sample.fit_resample(x_train, y_train)
```

```
[30]: sns.countplot(y_smote)
```

```
[30]: <AxesSubplot:xlabel='Class', ylabel='count'>
```



```
[33]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
# training
model.fit(x_smote, y_smote)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	71079
1	0.06	0.89	0.11	123
accuracy			0.98	71202
macro avg	0.53	0.93	0.55	71202
weighted avg	1.00	0.98	0.99	71202

F1 Score: 0.11202466598150052

```
[34]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_jobs=-1)
# training
model.fit(x_smote, y_smote)
# testing
```

```
# Determine number of fraud cases in dataset

df['Time'] = df['Time'] / 3600
fraud = df[df['Class'] == 1]
valid = df[df['Class'] == 0]

outlierFraction = len(fraud)/float(len(valid))

print(outlierFraction)
print('Fraud Cases: {}'.format(len(df[df['Class'] == 1])))


print('Valid Transactions: {}'.format(len(df[df['Class'] == 0])))
```

[49]:

```
0.0017304750013189597
Fraud Cases: 492
Valid Transactions: 284315
```

1.1 visualization

```
[50]: plt.figure(figsize=(12, 6))
plt.scatter(fraud['Time'], fraud['Amount'],
            color='red', marker='o', label='Fraud',
            alpha=0.6)
plt.scatter(valid['Time'], valid['Amount'], color='blue',
            marker='.', label='Non-Fraud', alpha=0.2)
plt.title('Time Series of Fraud vs. Non-Fraud Transactions')
plt.xlabel('Time (hours)')
plt.ylabel('Transaction Amount')
plt.legend(loc='upper right')
plt.show()
```



```
# dividing the X and the Y from the dataset
X = df.drop(['Class'], axis = 1)
Y = df["Class"]
print(X.shape)
print(Y.shape)
# getting just the values for the sake of processing
# (its a numpy array with no columns)
xData = X.values
yData = Y.values
```

(284807,)

ACCURACY

```
n_outliers = len(fraud)
n_errors = (yPred != yTest).sum()
print("The model used is Random Forest classifier ")
acc = accuracy_score(yTest, yPred)
print("The accuracy is {}".format(acc))
```

```
The model used is Random Forest classifier
The accuracy is 0.9994908886626171
So the Accuracy of the Model is 99.94
```

CONCLUSION:

- ▶ WE INVESTIGATED THE DATA, CHECKED FOR DATA UNBALANCING, VISUALIZED, AND UNDERSTOOD THE RELATIONSHIP BETWEEN DIFFERENT FEATURES. WE THEN USED FOUR PREDICTIVE MODELS TO PERFORM VALIDATION BY SPLITTING DATASET INTO 3 PARTS, A TRAIN SET, A VALIDATION SET AND A TEST SET. FOR THE FIRST TWO MODELS, WE ONLY USED THE TRAIN AND TEST SET.

THANK YOU